

Paper I — External Adversarial Audit — Final Report

External Adversarial Audit — Paper I (Split-Graph Fractional Triangle-Packing Theorem)

Audit ID: PAPER_I_EXT_AUDIT_2026-07-10

Date: 2026-07-10

Auditor: independent external AI adversarial auditor (adversarial-math-audit protocol)

Request spec: PAPER_I_EXTERNAL_FULL_AUDIT_REQUEST_SPEC.md

Target (only): Theorem 1.1 of *Affine Profile Reduction for Fractional Triangle Packings in Split Graphs* —

> For every split graph G on n vertices, $|E(G)| - 2 \cdot v_{\text{max}}^*(G) \leq n^2/6 + n$,
> where $v_{\text{max}}^*(G)$ is the maximum fractional triangle-packing value.

Ground truth: 01_manuscript/PAPER_I_preprint_v1.0.{md,pdf,tex} (Theorem 1.1, Appendix C); calculation ledger 02_validation/CALCULATION_LEDGER_v1.0.md (L1–L63); Lean sources 05_formalization/lean/.

Prior internal status: ledger FINAL_PASS / FROZEN / CLEAN; manuscript claims the theorem is machine-verified in Lean, sorry-free, axioms `[propext, Classical.choice, Quot.sound]`.

1. Executive verdict

Verdict: PASS (intermediate AI-assurance tier — not human peer review)

The finite fractional split-graph theorem $|E(G)| - 2v_{\text{max}}^*(G) \leq n^2/6 + n$ is **sound** and **correctly established** at every step the auditor could check independently, and the manuscript's scope is stated honestly (no over-claim into integral packing, clique partitions, chordal graphs, or Erdős #81). Specifically:

- **Analytic proof (Gates A–G):** the two-phase packing, the exact deficit identity, the affine/concavity reduction to pure profiles, the orbit symmetrization, the three-orbit LP, the quadratic benchmark $R(p, q)$, and the conversion to n all reproduce exactly. Every closed-form algebraic identity of the ledger was re-derived **symbolically** and **found exact** (Script 01, 9/9). The residual-cover LP optimum equals the closed form $\min\{U, D, H\}$ and the full $C(p, 2)$ -variable LP equals the 3-orbit LP (orbit symmetrization lossless), and the affine reduction with uniform bound $R(p, q) = p/2$

holds on 400 random mixed profiles with **no counterexample** (Script 02).

- **Falsification (Gate F / kill-switch):** on 622 split graphs (structured + random + adversarial) the theorem was **never violated**; the extremal family $K_p \vee K_{2p}$ attains $n^2/6 + n/6$ exactly, confirming the leading constant $1/6$ is tight and the $+n$ term is a non-optimized slack (Script 03).
- **Formal certificate (Gate H): PASS.** An **independent rebuild** in a fresh source tree (no inherited .lake cache; community Mathlib cache) reports `Build completed successfully (8027 jobs), exit 0, with PaperI.paperI\_main` **and** the LP-duality node `PaperI.Split.residual_duality` each depending on axioms `[propext, Classical.choice, Quot.sound]` — no `sorryAx`, no `native_decide`, no project-specific axiom. The released Lean sources match the recorded integrity manifest by SHA-256 (8/8), the Lean statement `paperI_main :  $\Phi \leq n^2/6 + n$  ( $\Phi = \text{edgeCount} - 2 \cdot \text{nu3star}$ ) is exactly Theorem 1.1, and a static scan finds no sorry/admit/native_decide/added axiom in the sources. This independently reproduces the recorded gate log.`
- **Citations (Gate I):** the single proof-critical import — finite-dimensional LP strong duality — is standard, and the Lean development derives it from Mathlib's finite Farkas (`residual_duality`), so the result does not depend on the exact Schrijver reference. No counterexample, circularity, unmet hypothesis, or scope over-claim was found. This is an independent AI adversarial audit: an intermediate assurance tier, stronger than self-audit but **not** external human peer review; it confers no publication status.

2. Scope, method, independence

In scope: Theorem 1.1 (finite fractional bound for split graphs) and its full proof chain, plus the Lean certificate for `paperI_main`.

Out of scope (confirmed absent from the claim): integral triangle packing, clique partitions/cp, chordal graphs, Φ_T , asymptotic Erdős #81, any fractional-to-integral transfer. Paper II and the asymptotic #81 route were not consulted.

Method: clean-context decomposition into claim IDs mapped to gates A–J; predeclared kill-switches for every computation; independent re-implementation of all primitives ($\vee$ LP, residual-cover LP, orbit LP, symbolic algebra); literal statement/definition cross-check across manuscript $\leftrightarrow$ ledger $\leftrightarrow$ Lean; independent Lean build in a fresh tree; SHA-256 provenance. **Independence:** fresh auditor code; fresh Lean source copy with no inherited .lake cache; the author's repo of record was not used. See `ENVIRONMENT.md`.

3. Gate table (A–J)

Gate	Item	Status	Evidence
A	Definitions (v^* , deficit, split model, κ , profiles, orbits) well-posed & consistent across manuscript/ledger/Lean	VERIFIED_WITHIN_SCOPE	manuscript §2 $\leftrightarrow$ ledger L1–L4 $\leftrightarrow$ Lean Split,Phi,nu3star,edgeCount,n cross-checked
B	Aggregate-load / affine profile reduction; candidates suffice	VERIFIED_WITHIN_SCOPE	Lemma 3.1 (L4) exact; concavity L33/L34; Script 02(B) 400 mixed profiles, 0 counterexample
C	Orbit symmetrization + finite LP; solved independently	VERIFIED_WITHIN_SCOPE	Script 02: full C(p,2)-LP = 3-orbit LP = $\min\{U, D, H\}$ (0/660 & 0/280, err 5e-15)
D	Finite LP duality applied within hypotheses; no infinite-dim appeal	VERIFIED_WITHIN_SCOPE	App B: finite feasible bounded pair; weak duality + attainment; Lean derives from Farkas
E	Deficit identity (4.9/L28) + quantitative bound $R(p, q)$; constant recomputed	VERIFIED_WITHIN_SCOPE	Script 01: L28, L48, L50, L51, L53, L58, q=0 all exact (9/9)
F	Conversion to n ($\leq n^2/6+n$) + comparison with optimum packing; tightness	VERIFIED_WITHIN_SCOPE	L58–L63 exact; Script 03: 0 violations/622, extremal attains $n^2/6+n/6$
G	Assembly composes to the theorem; no circularity	VERIFIED_WITHIN_SCOPE	linear chain L1→L63; Lean paperI_main assembles nu3star_ge_Vcom+Mcov_lower+edgeCount_eq+arithmetic
H	Lean: build, sorry-freeness, #print axioms, statement match, SHA	VERIFIED_WITHIN_SCOPE	Script 04: independent rebuild 8027 jobs exit 0; paperI_main+residual_duality axioms [proptext,Classical.choi ce,Quot.sound], no sorryAx; SHA 8/8 match; statement = Thm 1.1
I	Citations (LP duality; split-graph structure) applied within hypotheses	VERIFIED_WITH_RESIDUAL_RISK	finite LP duality standard + Lean-derived; exact "Cor. 7.1g" label not book-checked
J	Over-claim / scope check (no chordal, cp, #81, integral)	VERIFIED_WITHIN_SCOPE	manuscript §9, ledger §15 explicitly disclaim; statement is fractional split-graph only

4. Independent computation (Scripts 01–03)

All scripts declared kill-switches before running; seed 20260710; artifacts (script, results.txt, report.pdf, .zip, .sha256) under scripts/.

- **01 — algebraic identities (Gate E).** sympy exact check of L48, L50, L51 (SOS), L53, L58, the q=0 facts, and L25/L28. **9/9 exact** → CALCULATION_VERIFIED.
- **02 — orbit LP / symmetrization / reduction (Gates B, C).** (C0) q=0 branch verified; (C1) orbit LP = $\min\{U, D, H\}$ (0/660); (C2) full LP = orbit LP, err 5e-15 (0/280) — orbit averaging is lossless; (B) $M(\kappa) \geq \min_s M(\kappa^s)$ and $M(\kappa) \geq R-p/2$ on 400 mixed profiles, 0 violations, bound tight (slack 0). Auditor correction CORR-02-01 recorded

(a $q=0$ domain over-reach in the first run; not a paper defect).

- **03 — split-graph falsification (Gate F).** 622 split graphs, 0 theorem

violations; max ratio $(|E| - 2v_{\text{min}}^*) / (n^2/6 + n) = 0.815$; extremal $K_{p \vee K_{\lfloor 2p \rfloor}}$ attains $n^2/6 + n/6$ exactly (0 tightness failures).

5. Falsification / kill-switch log

No counterexample was found to Theorem 1.1 or to any gate. Failed falsification attempts (supporting evidence, not proof): adversarial dense complete-split graphs and 600 random split graphs (Script 03); random mixed neighborhood profiles probing the affine reduction and the uniform bound (Script 02); a search for an orbit/candidate beating $\min\{U, D, H\}$ (Script 02, none). The one issue surfaced was an **auditor** domain error (CORR-02-01), not a defect in the paper. Bounded search is existential evidence; the universal statement rests on the analytic proof + the Lean certificate.

6. Coverage declaration

Verified (within scope): all algebraic identities of the ledger (exact); the orbit LP closed form and the losslessness of orbit symmetrization (LP, bounded p); the affine reduction and uniform bound (LP, bounded p , random mixed profiles); the end-to-end inequality and extremal tightness (bounded split graphs); definition consistency across manuscript/ledger/Lean; Lean statement = Theorem 1.1; source SHA-256 = recorded manifest; static sorry/axiom-cleanliness of the sources; scope non-over-claim.

Residual risks: (i) numeric LP/enumeration checks are over bounded ranges ($p \leq 11$ for the LP identities, $p \leq 7-8$ for exact v_{min}^*), so they corroborate rather than prove the universal statement — but the analytic proof (Gates A–G) and the Lean certificate (Gate H) carry the universal claim; (ii) the exact wording of Schrijver "Corollary 7.1g" was not verified against the physical book (immaterial: the duality is standard and Lean-derived from Farkas); (iii) this is an AI audit, not human peer review.

7. Recommendation

Maintain the theorem at its internal FINAL_PASS / FROZEN status for the ****finite fractional split-graph bound only****. The independent audit found the analytic proof, the computational checks, and the formal certificate mutually consistent and free of any detected error, circularity, or over-claim. Human expert / journal peer review remains the appropriate next tier; this AI audit confers no publication status.

Distinctions maintained: "analytic proof correct" (Gates A–G, I), "Lean certificate holds" (Gate H), and "computation found no counterexample" (Scripts 02–03) are reported separately; all three point the same way for Paper I.